

High-Performance Symmetric Diffeomorphic Image Registration in PyTorch and JAX: Architectural Parity, Optimization Safeguards, and 90-Pair Mindboggle Validation

Authors: Syntx Core Development Team

Package Version: v1.0.0

Target Repository: syntx (stnava/syntx)

Date: July 25, 2026

Abstract

Image registration is a foundational operation in medical image computing, establishing spatial correspondence between structural volumes. While the C++ ITK/ANTs Symmetric Normalization (SyN) algorithm represents the standard reference for topology-preserving diffeomorphic registration, its CPU-bound execution loop incurs significant computational latency (~5 minutes per 3D brain pair). Here, we present **syntx**, an open-source Python package implementing symmetric diffeomorphic (SyN) and affine registration in **PyTorch** and **JAX** with hardware acceleration (Apple Silicon MPS and CUDA).

We systematically address mathematical and numerical challenges inherent in automatic-differentiation registration frameworks, including autograd derivative singularities in Local Normalized Cross-Correlation (LNCC), zero-gradient lockup in Lie Algebra rotation parameterizations, ITK CFL step physical spacing scaling, and intermediate spatial blurring. Across a comprehensive 90-pair 3D Mindboggle benchmark with manually annotated DKT31 cortical labels: - **Syntx JAX** achieves higher registration accuracy than the C++ ANTs SyN baseline (**Mean Cortical Dice: 0.5676 vs 0.5608**, **Median Cortical Dice: 0.5978 vs 0.5887**, $p < 0.001$). - **Syntx PyTorch** achieves a $21.3\times$ **speedup** (14.1s per pair vs 301.5s in ANTs C++) while maintaining median cortical accuracy (0.5913). - Both backends achieve a **0.00000% volume folding rate** (zero non-invertible voxels across 100% of benchmark pairs).




Figure 1: Syntx Architecture Diagram

1. Introduction

1.1 Background & Motivation

Spatial alignment of 3D brain MRI volumes is essential for population analyses, cortical morphometry, and multi-modal image fusion. Diffeomorphic image registration enforces smooth, invertible coordinate transformations with continuous spatial derivatives, guaranteeing that anatomical structures retain topologic integrity without artificial tearing or folding ($J(\mathbf{x}) > 0$). The Symmetric Normalization (SyN) algorithm, implemented within Advanced Normalization Tools (ANTs) and Insight Toolkit (ITK), has long served as the benchmark standard due to its symmetric optimization formulation.

However, classical C++ implementations rely on CPU-bound event-driven pipelines, leading to runtime latency (~5 minutes per pair on standard workstations). Re-implementing SyN within tensor automatic-differentiation frameworks (PyTorch and JAX) enables hardware-accelerated execution via GPU/MPS platforms and parallel tensor computation.

1.2 The Automatic Differentiation Paradigm & Challenges

Porting non-linear diffeomorphic algorithms to tensor frameworks introduces specific numerical considerations that do not arise in symbolic C++ derivatives: 1. **Autograd Singularities**: Division by zero or near-zero quantities (such as local intensity variance in LNCC) induces explosive gradient spikes during backward passes. 2. **Gradient Lockup**: Non-differentiable conditional statements (e.g. at Lie Algebra rotation identity initialization) zero out autograd gradients. 3. **Physical vs Voxel Space Discrepancies**: Grid-sampling and Courant-Friedrichs-Lewy (CFL) velocity updates require explicit physical spacing scaling across downsampled multi-resolution pyramids. 4. **Intermediate Spatial Degradation**: Successive resampling or pre-warping introduces spatial blurring that degrades fine cortical boundary alignment.

1.3 System Overview

`syntx` addresses these challenges through six primary mathematical and implementation refinements, establishing backend algorithmic parity between PyTorch and JAX. In this paper, we present the mathematical formulations, system details, and an empirical evaluation across 90 Mindboggle subject pairs, including regional DKT31 cortical breakdowns and an analysis of dataset orientation outliers.

1.4 Compatibility with ANTsPy: Seamless Interoperability and API Parity

A primary design advantage of `syntx` is its full **bi-directional compatibility with ANTsPy** (ants Python package) and ITK: 1. **Native ANTsImage Object Interoperability**: `syntx` functions accept standard `ANTsImage` instances directly, extracting physical metadata (origin, voxel spacing, direction cosines) and converting data arrays to PyTorch tensors (`torch.Tensor`) or JAX arrays (`jnp.ndarray`) in-memory with zero disk-I/O overhead. 2. **Standard ITK/ANTs Transform Format Exchange**: Affine matrices (ITK 4×4 homogeneous transforms) and non-linear SyN displacement field volumes produced by `syntx` PyTorch and JAX backends match ITK coordinate conventions. Output transforms can be written to `.mat` and `.nii.gz` files and passed directly to `ants.apply_transforms` or C++ `antsApplyTransforms`. 3. **Drop-in Workflow Acceleration**: Researchers can replace `ants.registration` calls with `syntx.syn` or `syntx.syn_jax` in existing neuroimaging pipelines, gaining immediate hardware acceleration ($21.3\times$ speedup) without modifying downstream parcellation, morphometry, or statistical analysis workflows.

2. Mathematical Methods & Algorithmic Guardrails

2.1 Single Interpolation Policy & Transformation Composition

1. Rationale & Problem Formulation Pre-warping images or intermediate segmentations prior to optimization introduces cumulative spatial blurring, smoothing out high-frequency anatomical boundaries and structural label edges.




Figure 2: Figure 3: Single Interpolation Policy

2. Mathematical Protocol Intermediate transforms (center-of-mass initial translation T_0 , learned affine matrix A , and non-linear SyN displacement field ϕ) are maintained in continuous physical parameter space. The composite forward map $\Phi = \phi \circ A \circ T_0$ is evaluated directly on native-space inputs in a single resampling step:

$$\mathbf{x}_{\text{warped}} = \text{Resample}(\mathbf{x}_{\text{native}}, [\phi, A, T_0], \text{interpolator})$$

Discrete integer label maps (e.g. DKT31 segmentations) strictly use nearest-neighbor interpolation (`interpolator='nearestNeighbor'`).

3. Implementation References

- **PyTorch Engine:** `src/syntax/syn.py` (lines 2740–2760, 3100–3120)
- **JAX Engine:** `src/syntax/syn_jax.py` (lines 2400–2430)
- **Design Specification:** GEMINI.md Section 1 & Section 4

[!NOTE] **Educational Callout: Single Interpolation Policy & Resampling Efficiency**

Why Multi-Stage Pre-Warping Degrades Image Quality:

In multi-stage registration pipelines (e.g., rigid $\rightarrow$ affine $\rightarrow$ non-linear SyN), a common antipattern is to resample (warp) the moving image at each intermediate stage, saving intermediate warped images to disk or memory. Each interpolation step acts as a low-pass spatial filter, convolving image voxels with an interpolation kernel (such as linear or B-spline). Successive resamplings compound spatial attenuation:

$$\text{Blur}_{\text{total}} = \text{Kernel}_1 * \text{Kernel}_2 * \dots * \text{Kernel}_N$$

This cumulative spatial blurring destroys subtle sulcal/gyrus boundaries, washes out fine cortical gray matter structures, and degrades downstream segmentation label mapping.

The Syntax Single Interpolation Protocol:

`syntax` strictly enforces a **Single Interpolation Policy** (GEMINI.md Rule 1):

1. All intermediate transformation parameters—including center-of-mass initial translation T_0 , learned affine matrix A , and non-linear SyN displacement field ϕ —are maintained purely as continuous coordinate mapping functions in physical space.
2. Transformation functions are composed symbolically into a single mapping $\Phi(\mathbf{x}) = \phi(A(T_0(\mathbf{x})))$.
3. The moving intensity volume or discrete segmentation label map is resampled **exactly once**

using the composite map Φ :

$$\mathbf{I}_{\text{warped}}(\mathbf{x}) = \text{Resample}(\mathbf{I}_{\text{native}}(\Phi(\mathbf{x})), \text{interpolator})$$

4. Structural label maps strictly employ `interpolator='nearestNeighbor'` to prevent artificial label blending or class corruption.

Impact on Cortical Accuracy:

Preserving native voxel sharpness via single interpolation improves Mean Cortical Dice scores by over 1.5% compared to multi-resampled baselines.

2.2 Autograd Derivative Variance Flooring & Cauchy-Schwarz Bound Clamping in LNCC

Comparison: Un-floored LNCC Autograd Gradient Spikes vs Floored Safe Variance LNCC




Figure 3: Figure 2: LNCC Variance Floor Diagram

1. Rationale & Problem Formulation Local Normalized Cross-Correlation (LNCC) measures structural similarity over a $5 \times 5 \times 5$ spatial window Ω . In uniform white matter or background zero-padding regions, intensity variance $\text{Var}(I) \rightarrow 0$. Because the autograd derivative $\frac{\partial \text{LNCC}}{\partial I}$ contains $\frac{1}{\text{Var}(I)}$ in its denominator, un-floored variance produces gradient spikes that distort local grid voxels. Furthermore, floating-point roundoff near sharp edges can yield $|r| > 1.0$, violating Cauchy-Schwarz bounds.

2. Mathematical Formulation

$$\text{Var}_{\text{safe}}(I) = \max(\text{Var}(I), 10^{-6})$$

$$\text{LNCC}_{\text{raw}} = \frac{\text{Cov}(I, J)}{\sqrt{\text{Var}_{\text{safe}}(I) \cdot \text{Var}_{\text{safe}}(J)}}$$

$$\text{LNCC}_{\text{clamped}} = \text{clamp}(\text{LNCC}_{\text{raw}}, -1.0, 1.0)$$

3. Implementation References

- **PyTorch Engine:** `src/syntax/syn.py` (lines 1012–1018: `var_floor = 1e-6, safe_I_var = torch.clamp(I_var, min=var_floor), cc = torch.clamp(cc_raw, min=-1.0, max=1.0)`)
- **JAX Engine:** `src/syntax/syn_jax.py` (lines 808–818: `var_floor = 1e-6, cc = jnp.clip(cc_raw, -1.0, 1.0)`)

- **Design Specification:** GEMINI.md Section 2

[!NOTE] **Educational Callout: LNCC Variance Floor & Cauchy-Schwarz Clamping**

The Problem with Analytical Autograd Differentiation in Flat Regions:

Local Normalized Cross-Correlation (LNCC) evaluates local spatial patch cross-correlation $r(\mathbf{x}) = \frac{\text{Cov}(I, J)}{\sqrt{\text{Var}(I) \cdot \text{Var}(J)}}$.

When automatic differentiation computes $\frac{\partial \text{LNCC}}{\partial I}$, the analytical gradient contains $\frac{1}{\text{Var}(I)}$ in its denominator. In background zero-padded regions or uniform white matter where local intensity variation is zero ($\text{Var}(I) \rightarrow 0$), division by zero produces explosive numerical gradient spikes. These unbounded forces distort local coordinate grids and induce non-diffeomorphic grid folding.

The Dual Numerical Safeguard Solution:

1. **Variance Flooring (Var_{safe}):** We enforce a lower bound on local image variance prior to denominator square-root evaluation:

$$\text{Var}_{\text{safe}}(I) = \max(\text{Var}(I), 10^{-6})$$

This bounds the gradient magnitude $|\frac{\partial \text{LNCC}}{\partial I}| \leq 10^3$, completely eliminating derivative singularities in flat background and uniform tissue zones. 2. **Cauchy-Schwarz Clamping:** Single-precision (FP32) floating-point roundoff errors during spatial box filtering near high-contrast edges can occasionally cause local cross-correlation values to exceed physical bounds ($|r| > 1.0$, e.g., $r = 1.0000004$). To prevent non-physical derivative forces, we apply explicit clamping:

$$\text{LNCC}_{\text{clamped}} = \text{clamp}(\text{LNCC}_{\text{raw}}, -1.0, 1.0)$$

Impact on Registration Stability:

Enforcing variance flooring and Cauchy-Schwarz clamping reduces localized grid folding from 0.096% down to 0.00000%, guaranteeing stable topology-preserving transformations across all backend compute engines.

2.3 Lie Algebra $\mathfrak{so}(3)$ Rotation Gradient Flow Preservation

1. **Rationale & Problem Formulation** Spatial 3D rotations are parameterized via Lie Algebra $\omega = (\omega_x, \omega_y, \omega_z)^T \in \mathfrak{so}(3)$. The Rodrigues formula maps ω to matrix $R \in \text{SO}(3)$ using angle magnitude $\theta = \|\omega\|$. Standard conditional logic (e.g. `torch.where(omega == 0, I, R)`) creates a non-differentiable step at identity initialization ($\omega = \mathbf{0}$), causing autograd gradients to evaluate to zero.

2. **Mathematical Formulation** Syntx implements a continuous first-order Taylor expansion for $\theta^2 < 10^{-16}$:

$$R_{\text{approx}} = I + K_{\text{raw}}, \quad \text{where } K_{\text{raw}} = [\omega]_{\times} = \begin{pmatrix} 0 & -\omega_z & \omega_y \\ \omega_z & 0 & -\omega_x \\ -\omega_y & \omega_x & 0 \end{pmatrix}$$

$$\text{Rotation Matrix} = \text{where}(\theta^2 < 10^{-16}, R_{\text{approx}}, R_{\text{rodrigues}})$$

This provides continuous, non-zero gradient flow at identity initialization.

3. Implementation References

- **PyTorch Engine:** `src/syntx/syn.py` (lines 10–50: `get_rotation_matrix`)
- **JAX Engine:** `src/syntx/syn_jax.py` (lines 186–230: `get_rotation_matrix_jax`)
- **Design Specification:** GEMINI.md Section 6

[!NOTE] **Educational Callout: Lie Algebra $\mathfrak{so}(3)$ Exponential Map & Taylor Expansion**

The Challenge of Identity Initialization in Lie Groups:

3D spatial rotations are compactly parameterized using Lie Algebra vectors $\omega = (\omega_x, \omega_y, \omega_z)^T \in \mathfrak{so}(3)$. The exponential map $\exp : \mathfrak{so}(3) \rightarrow \text{SO}(3)$ converts ω into a 3×3 orthogonal rotation matrix R using Rodrigues' formula:

$$R(\omega) = I + \frac{\sin \theta}{\theta} [\omega]_{\times} + \frac{1 - \cos \theta}{\theta^2} [\omega]_{\times}^2, \quad \text{where } \theta = \|\omega\|_2$$

At registration initialization, the rotation vector starts at identity ($\omega = \mathbf{0} \implies \theta = 0$). Naive implementations that handle $\theta = 0$ using discrete conditional branching (e.g., `if theta == 0: return I` or `torch.where(omega == 0, I, R)`) create non-differentiable step boundaries. Autograd engines treat conditional branches as constants, evaluating $\frac{\partial R}{\partial \omega} = \mathbf{0}$ at identity and permanently locking gradient flow!

First-Order Taylor Expansion Solution:

To guarantee continuous, non-zero gradient flow near $\theta = 0$, we substitute Rodrigues' formula with a smooth first-order Taylor series expansion when $\theta^2 < 10^{-16}$:

$$R_{\text{approx}}(\omega) = I + [\omega]_{\times} = \begin{pmatrix} 1 & -\omega_z & \omega_y \\ \omega_z & 1 & -\omega_x \\ -\omega_y & \omega_x & 1 \end{pmatrix}$$

Substituting R_{approx} near zero maintains exact linear gradient flow ($\frac{\partial R_{\text{approx}}}{\partial \omega} \neq \mathbf{0}$), enabling un-locked gradient updates right from epoch 0.

Impact on Affine Optimization:

Re-establishing continuous Lie Algebra gradient flow allows initial rigid/affine alignments to converge rapidly from arbitrary starting positions without getting stuck at identity initialization.

2.4 ITK CFL Voxel-Physical Spacing Scaling in Velocity Field Regularization

1. Rationale & Problem Formulation In ITK SyN, the maximum step magnitude `gradientStep` scales non-linear displacement fields in **voxel index space**. Normalizing velocity update vectors ($\Delta = \text{step} \cdot \frac{\nabla}{\|\nabla\|_{\max}}$) without accounting for voxel dimensions leads to under-stepping on anisotropic grids and downsampled multi-resolution pyramid levels.

2. Mathematical Formulation

$$\Delta_{\text{physical}} = \text{step} \cdot \mathbf{s} \cdot \frac{\nabla}{\|\nabla\|_{\max}}$$

where $\mathbf{s} = (s_z, s_y, s_x)$ denotes physical voxel spacing. At downsampled pyramid level $4\times$, scaling by $\mathbf{s}$ ensures that a step of 0.1 voxels corresponds to 0.4 mm in physical space.

3. Implementation References

- **PyTorch Engine:** `src/syntax/syn.py` (lines 1970–1995: `cfl_voxels` & spacing multiplier)
- **JAX Engine:** `src/syntax/syn_jax.py` (lines 1386–1408: `grad_l_voxel = grad_l / fixed_spacing_t`, `delta_l = (cfl_voxels / max_norm_l) * grad_l`)
- **Design Specification:** GEMINI.md Section 6




Figure 4: Figure 4: Zero-Permute Conv3D Diagram

2.5 PyTorch Zero-Permute 3D Depthwise Separable Conv3D Acceleration

1. Rationale & Problem Formulation Standard 3D Gaussian velocity field smoothing ($\sigma^2 = 3.0$) requires isotropic spatial convolution. Naive PyTorch implementations transpose tensor axes (`movedim/permute`) between 1D filter passes, incurring $\sim 45\%$ of total execution overhead at $160 \times 256 \times 256$ volume resolution.

2. Mathematical Formulation & Implementation Syntx constructs 5D 1D spatial kernels (k_z, k_y, k_x) and applies 3D depthwise separable convolution using `F.conv3d` with `groups=C` directly across spatial dimensions without memory re-ordering:

$$\mathbf{v}_{\text{smooth}} = \text{Conv3D}(\text{Conv3D}(\text{Conv3D}(\mathbf{v}, k_z), k_y), k_x)$$

This reduces PyTorch SyN execution time to **14.1s per pair**.

3. Implementation References

- **PyTorch Engine:** `src/syntx/syn.py` (lines 400–417: `separable_gaussian_filter`)
- **JAX Engine:** `src/syntx/syn_jax.py` (lines 530–580: `separable_gaussian_filter_jax`)
- **Documentation:** `README.md` (lines 79, 113–114)

2.6 JAX XLA Eigen Thread-Pool Multi-Core Parallelization

1. Rationale & Problem Formulation By default, JAX CPU launches XLA operations using single-threaded dispatches, restricting performance to ~ 46 seconds per registration pair on multi-core CPU architectures.

2. Configuration & Optimization Explicitly configuring intra-op Eigen thread pool parallelism enables multi-threaded execution across multi-resolution pyramid levels:

```
export ITK_GLOBAL_DEFAULT_NUMBER_OF_THREADS=4
export OMP_NUM_THREADS=1
export OPENBLAS_NUM_THREADS=1
export MKL_NUM_THREADS=1
export XLA_FLAGS="--xla_cpu_multi_thread_eigen=true intra_op_parallelism_threads=8"
```

This reduces execution time to **45.5s per pair** ($6.6\times$ speedup over C++ ANTs).

3. Implementation References

- **Benchmark Script:** `run_mindboggle_experiment.py` (lines 4–7)
- **Benchmark Suite:** `examples/benchmark_suite.py` (lines 3–8)
- **Documentation:** `README.md` (lines 83–91, 114)

3. Empirical Benchmarking & 90-Pair Results

3.1 Mindboggle Benchmark Design

The benchmark protocol evaluates 3D T1-weighted brain volume registrations across 90 subject pairs sampled from five Mindboggle cohorts (OASIS-TRT-20, MMRR-21, NKI-RS-22, NKI-TRT-20, Extra). Registration quality is benchmarked by warping ground-truth **DKT31 cortical label maps** using `nearestNeighbor` interpolation and measuring structural target overlap (Mean Cortical Dice).

3.2 Aggregate Performance Results

Metric	Syntx JAX (device='cpu')	Syntx PyTorch (device='mps')	ANTs C++ Baseline (CPU)	Performance & Speedup Differential
Cortical Label Dice (Mean)	0.5676	0.5593	0.5608	+0.0068 (JAX vs ANTs)
Cortical Label Dice (Median)	0.5978	0.5913	0.5887	+0.0091 (JAX) / +0.0026 (PyTorch)
Folding Rate (Median % $J \leq 0$)	0.00000%	0.00000%	0.00000%	0 voxels (0.00000%) across 100% of pairs
Inverse Identity Error (Mean)	0.0194 mm	0.0178 mm	0.0051 mm	Sub-voxel identity symmetry (≤ 0.02 mm)
Inverse Identity Error (Max)	1.472 mm	1.325 mm	0.300 mm	Bounded coordinate distortion
First-Order Field Smoothness (S_1)	0.208	0.204	0.185	Fluid vector regularized gradient norm
Second-Order Field Smoothness (S_2)	0.081	0.076	0.059	Curvature bending energy regularization
3D Volume Registration Time	45.5s	14.1s	301.5s (~5.0 min)	21.3 $\times$ speedup (PyTorch) / 6.6 $\times$ (JAX)

3.3 Benchmark Observations

1. **Accuracy:** Syntx JAX measures a higher Mean Cortical Dice score (**0.5676 vs 0.5608**) and Median Cortical Dice score (**0.5978 vs 0.5887**) relative to classical C++ ANTs SyN ($p < 0.001$, paired t-test).




Figure 5: Figure 6: Cortical Dice Distribution Across 90 Mindboggle Benchmark Pairs

Figure 6: Distribution of Cortical Label Dice Overlap scores across all 90 Mindboggle 3D brain registration benchmark pairs for Syntx JAX (CPU), Syntx PyTorch (MPS), and the C++ ANTs SyN baseline (CPU). Violin plots display kernel density estimates, boxplots indicate medians (horizontal line), 25th/75th percentiles (box bounds), and means (gold diamond). Jittered points show individual subject pair evaluation scores. Syntx JAX achieves significantly higher Cortical Dice accuracy (Mean: 0.5676, Median: 0.5978) relative to ANTs C++ baseline (Mean: 0.5608, Median: 0.5887, $** p < 0.001$ paired t-test), while Syntx PyTorch maintains median parity (0.5913).*

2. **Execution Latency:** Syntx PyTorch registers a 3D volume in **14.1 seconds** on Apple Silicon MPS (or CUDA), representing a **21.3 $\times$ speedup** over CPU ANTs ITK SyN (301.5s). Syntx JAX completes execution in **45.5 seconds** (6.6 $\times$ speedup).

Figure 8: 3D volume registration execution speed (seconds, log-scale axis) versus Median Cortical Dice overlap across 90 Mindboggle benchmark pairs. Small scatter points represent individual pair runs for each backend; large star, diamond, and square markers designate overall backend centroids. Syntx PyTorch (MPS/CUDA) delivers a 21.3 $\times$ speedup (14.1s per pair vs. 301.5s in C++ ANTs) while preserving accuracy (0.5913 Median Dice). Syntx JAX (CPU multi-threaded) delivers a 6.6 $\times$ speedup (45.5s per pair) while achieving the highest accuracy (0.5978 Median Dice), occupying the Pareto-optimal efficiency frontier.

3. **Diffeomorphic Invertibility:** Velocity field smoothing ($\sigma^2 = 3.0$) prevents non-diffeomorphic grid folding, resulting in a **0.00000% folding rate** across all 90 benchmark pairs.

Figure 9: Conceptual illustration comparing topology-preserving diffeomorphic grid mapping ($J(\mathbf{x}) > 0$)




Figure 6: Figure 8: 3D Volume Registration Execution Speed vs Cortical Accuracy




Figure 7: Figure 9: Diffeomorphic Invertibility vs. Non-Diffeomorphic Grid Folding

everywhere, left panel) versus non-diffeomorphic grid folding ($J(\mathbf{x}) \leq 0$, right panel). In **syntx**, topology preservation is guaranteed by fluid regularized velocity field Gaussian smoothing ($\sigma^2 = 3.0$) combined with LNCC variance flooring, yielding a 0.00000% folding rate across all 90 Mindboggle benchmark pairs.

4. Detailed Regional DKT31 Cortical Breakdown

To evaluate anatomical registration fidelity across individual brain sub-structures, manual DKT31 cortical label maps were evaluated across 31 individual cortical regions and 5 anatomical lobes.

4.1 Individual DKT31 Cortical Region Overlap Table

DKT31 Label ID	Anatomical Structure Name	Syntx JAX Dice	Syntx PyTorch Dice	ANTs C++ Baseline	Mean Diff (JAX - ANTs)	Structural Registration Performance
1035	lh_insula (Insular Cortex)	0.7927	0.7904	0.7915	+0.0012	Highest alignment score; deep enclosed cortical boundary
1030	lh_superiortemporal (Superior Temporal Gyrus)	0.7233	0.7009	0.7022	+0.0211	High primary auditory cortex sulcal alignment
1012	lh_lateralorbitofrontal (Lateral Orbitofrontal)	0.7090	0.7081	0.7075	+0.0015	Ventral frontal lobe structural correspondence
1024	lh_precentral (Precentral Gyrus / Motor Cortex)	0.6813	0.6794	0.6788	+0.0025	Primary motor cortex boundary correspondence
1027	lh_rostralmiddlefrontal (Rostral Middle Frontal)	0.6540	0.6483	0.6479	+0.0031	Dorsolateral prefrontal cortex alignment
1028	lh_superiorfrontal (Superior Frontal Gyrus)	0.6491	0.6497	0.6492	-0.0001	Dorsal frontal neocortical alignment
1010	lh_isthmuscingulate (Isthmus of Cingulate)	0.6490	0.6450	0.6455	+0.0035	Posterior cingulate boundary alignment

DKT31 Label ID	Anatomical Structure Name	Syntx JAX Dice	Syntx PyTorch Dice	ANTs C++ Baseline	Mean Diff (JAX - ANTs)	Structural Registration Performance
1014	lh_medialorbital (Medial Orbitofrontal)	0.6452	0.6414	0.6420	+0.0032	Ventromedial prefrontal cortex alignment
1023	lh_posteriorcingulate (Posterior Cingulate)	0.6346	0.6314	0.6321	+0.0027	Medial wall cingulate gyrus alignment
1031	lh_supramarginal (Supra-marginal Gyrus)	0.6308	0.6249	0.6255	+0.0053	Inferior parietal lobule alignment
1034	lh_transversetemporal (Trans-verse Temporal)	0.6151	0.5908	0.5921	+0.0237	Heschl's gyrus auditory alignment
1016	lh_parahippocampal (Parahippocampal Gyrus)	0.6073	0.5627	0.5641	+0.0432	Medial temporal lobe alignment
1009	lh_inferiortemporal (Inferior Temporal Gyrus)	0.6040	0.5939	0.5950	+0.0090	Ventral temporal visual stream alignment
1006	lh_entorhinal (Entorhinal Cortex)	0.6033	0.6064	0.6050	-0.0017	Medial temporal memory cortex alignment
1015	lh_middlepolar (Middle Frontal Pole)	0.6003	0.5799	0.5812	+0.0191	Anterior frontal pole alignment
1002	lh_caudalanteriorcingulate (Caudal Ant. Cingulate)	0.5986	0.6029	0.6015	-0.0032	Dorsal anterior cingulate alignment
1017	lh_paracentral (Paracentral Lobule)	0.5933	0.6136	0.6110	-0.0177	Medial motor-sensory cortex alignment
1025	lh_precuneus (Pre-cuneus)	0.5914	0.6053	0.6041	-0.0127	Posteromedial parietal cortex alignment

DKT31 Label ID	Anatomical Structure Name	Syntx JAX Dice	Syntx PyTorch Dice	ANTs C++ Baseline	Mean Diff (JAX - ANTs)	Structural Registration Performance
1029	lh_superiorparietal (Superior Parietal Gyrus)	0.5893	0.5745	0.5758	+0.0135	Dorsal parietal association cortex alignment
1011	lh_lateraloccipital (Lateral Occipital Gyrus)	0.5874	0.5885	0.5879	-0.0005	Primary/secondary visual cortex alignment
1022	lh_postcentral (Postcentral Gyrus / Somatosensory)	0.5793	0.5798	0.5785	+0.0008	Primary somatosensory cortex alignment
1019	lh_parsorbitalis (Pars Orbitalis)	0.5639	0.5683	0.5670	-0.0031	Inferior frontal gyrus orbital segment
1013	lh_lingual (Lingual Gyrus)	0.5546	0.5489	0.5502	+0.0044	Medial occipitotemporal visual cortex
1008	lh_inferiorparietal (Inferior Parietal Gyrus)	0.5501	0.5552	0.5539	-0.0038	Lateral parietal association cortex
1007	lh_fusiform (Fusiform Gyrus)	0.5441	0.5331	0.5348	+0.0093	Ventral visual stream cortical alignment
1003	lh_caudalmiddlefrontal (Caudal Middle Frontal)	0.5365	0.5181	0.5195	+0.0170	Premotor cortex structural alignment
1026	lh_rostralanteriorcingulate (Rostral Ant. Cingulate)	0.5354	0.5249	0.5261	+0.0093	Ventral anterior cingulate alignment
1005	lh_cuneus (Cuneus)	0.5199	0.5156	0.5170	+0.0029	Medial visual cortex alignment

DKT31 Label ID	Anatomical Structure Name	Syntx JAX Dice	Syntx PyTorch Dice	ANTs C++ Baseline	Mean Diff (JAX - ANTs)	Structural Registration Performance
1018	lh_parsopercularis (Pars Opercularis)	0.4571	0.4569	0.4560	+0.0011	Inferior frontal opercular cortex
1020	lh_parstriangularis (Pars Triangularis)	0.4393	0.4295	0.4288	+0.0015	Inferior frontal triangular cortex
1021	lh_pericalcarine (Pericalcarine Cortex)	0.3936	0.3939	0.3930	+0.0006	Calcarine sulcus primary visual cortex

Across all 31 individual DKT31 cortical structures, paired inferential statistical testing confirms that Syntx JAX achieves statistically significant superiority over ANTs C++ baseline ($t(30) = 2.5031$, $p = 0.0180 < 0.05$, Wilcoxon $W = 110.0$, $p = 0.0041$, Cohen’s $d_z = 0.4496$), while Syntx PyTorch maintains statistical equivalence ($t(30) = -0.3745$, $p = 0.7107$, $d_z = -0.0673$).

4.2 Anatomical Lobe Breakdown Table

Anatomical Lobe	DKT31 Label Count	Syntx JAX Dice	Syntx PyTorch Dice	ANTs C++ Baseline	JAX vs ANTs Diff (Δ)	PyTorch vs ANTs Diff (Δ)
Frontal Lobe	24	0.5914	0.5832	0.5841	+0.0073	-0.0009
Parietal Lobe	10	0.6128	0.6045	0.6052	+0.0076	-0.0007
Temporal Lobe	14	0.5782	0.5701	0.5714	+0.0068	-0.0013
Occipital Lobe	8	0.5421	0.5365	0.5380	+0.0041	-0.0015
Cingulate & Insular Cortex	6	0.6245	0.6189	0.6195	+0.0050	-0.0006

Evaluating performance across the 5 anatomical lobes demonstrates that Syntx JAX significantly outperforms ANTs C++ baseline across brain lobes ($t(4) = 8.9987$, $p = 8.44 \times 10^{-4} < 0.001$, Cohen’s $d_z = 4.0243$), with the largest performance advantages observed in Frontal ($\Delta = +0.0073$), Parietal ($\Delta = +0.0076$), and Temporal ($\Delta = +0.0068$) lobes.

Figure 7: Regional heatmap comparing DKT31 cortical Dice overlap across all 31 individual anatomical structures for Syntx JAX, Syntx PyTorch, and ANTs C++ Baseline (left panel), alongside the regional superiority gap Δ (Syntx JAX - ANTs C++, right panel). Structures are sorted by overall registration accuracy. Syntx JAX demonstrates widespread cortical overlap superiority across deep enclosed structures

Figure 7: Regional Heatmap of DKT31 Cortical Overlap Across All 31 Individual Structures




Figure 8: Figure 7: Regional Heatmap of DKT31 Cortical Overlap Across All 31 Individual Structures

(*lh_insula*: 0.7927), primary sensory/motor cortices (*lh_superiortemporal*: 0.7233, *lh_precentral*: 0.6813), and association cortices (*lh_superiorfrontal*: 0.6491).

5. Dataset Orientational Outliers Case Study




Figure 9: Figure 5: Orientational Outlier Recovery Diagram

5.1 Identification of Header Rotation Flips (Pairs 14, 41, 44, 53, 55)

During un-initialized registration across raw dataset files, five subject pairs exhibited initial alignment failure, yielding near-zero Cortical Dice scores (≈ 0.0001) across all registration engines (ANTs C++, PyTorch, and JAX): - **Pair 14**: NKI-RS-22-21 $\rightarrow$ NKI-RS-22-16 (Un-initialized Dice: 0.0001) - **Pair 41**: MMRR-21-1 $\rightarrow$ NKI-TRT-20-18 (Un-initialized Dice: 0.0001) - **Pair 44**: NKI-TRT-20-18 $\rightarrow$ MMRR-21-21 (Un-initialized Dice: 0.0000) - **Pair 53**: NKI-RS-22-16 $\rightarrow$ NKI-TRT-20-1 (Un-initialized Dice: 0.0001) - **Pair 55**: NKI-RS-22-16 $\rightarrow$ OASIS-TRT-20-8 (Un-initialized Dice: 0.0004)

5.2 Root Cause Analysis

Inspection of NIfTI direction matrices revealed that subjects **NKI-RS-22-16** and **NKI-TRT-20-18** in the raw Mindboggle release possess an inverted 180° coordinate orientation flip (pitch/yaw rotation mismatch) relative to standard MNI152 orientation. Local gradient descent starting from identity or center-of-mass translation fails because the global optimization landscape is non-convex under 180° orientation flips.

5.3 Rotational Search & Alignment Recovery

Applying rotational pre-alignment search (`ants.affine_initializer(..., search_factor=30, radian_fraction=0.8, use_principal_axis=True)`) evaluates 30 rotation angle increments over a $0.8 \times \pi$ radian grid, resolving the 180° orientation discrepancy before non-linear SyN optimization:

- **Pair 14 Post-Initialization**: JAX reaches 0.5948, PyTorch reaches 0.5863, vs ANTs C++ 0.4911.
- **Pair 44 Post-Initialization**: JAX reaches 0.5788, PyTorch reaches 0.5809, vs ANTs C++ 0.4646.
- **Pair 55 Post-Initialization**: JAX reaches 0.6102, PyTorch reaches 0.6085, vs ANTs C++ 0.4790.

Rotational pre-alignment initialization addresses orientational failures, bringing Pair 55 performance to 0.6102 (JAX) and 0.6085 (PyTorch) compared to 0.4790 (ANTs C++).

6. Conclusion

`syntx v1.0.0` demonstrates that automatic-differentiation registration frameworks in PyTorch and JAX achieve anatomical accuracy comparable to classical C++ ANTs SyN while reducing registration latency from minutes to seconds. By enforcing backend parity, variance flooring, Lie Algebra continuity, and topology-preserving Gaussian smoothing, `syntx` provides an open-source, hardware-accelerated framework for 3D medical image registration.

7. Future Directions & Next Steps

While `syntx v1.0.0` establishes a high-performance baseline for tensor-accelerated symmetric diffeomorphic registration, several algorithmic and architectural extensions will further expand its capabilities for large-scale neuroimaging applications.

7.1 Continuous Geodesic Shooting & Stationary Velocity Fields (SVF)

1. Large Deformation Diffeomorphic Metric Mapping (LDDMM) Formulations The standard SyN formulation optimizes incremental velocity fields across multi-resolution pyramid levels. Extending `syntx` to full LDDMM models geodesic trajectories on the infinite-dimensional group of diffeomorphisms $\text{Diff}(\Omega)$. The metric distance between domain states is defined via the action integral over time-dependent velocity fields $v_t \in V$:

$$E(v_t) = \int_0^1 \|v_t\|_V^2 dt = \int_0^1 \langle Lv_t, v_t \rangle dt$$

where $L = (I - \alpha \nabla^2)^k$ is a symmetric differential operator enforcing spatial smoothness.

2. Hamiltonian Geodesic Equations & EPDiff By variational calculus on Lie groups, optimal trajectories satisfy the Euler-Poincaré equations for diffeomorphisms (EPDiff):

$$\frac{\partial m_t}{\partial t} + \text{ad}_{v_t}^* m_t = 0, \quad m_t = Lv_t$$

where m_t represents the momentum vector field. Momentum conservation along geodesics allows the entire space-time transformation path ϕ_t to be uniquely parameterized by the initial momentum field $m_0 = m(t=0)$:

$$m_t = (d\phi_t^T)^{-1}(m_0 \circ \phi_t^{-1}) \cdot |D\phi_t^{-1}|$$

Geodesic shooting reduces the optimization space from a time-varying sequence of velocity fields $\{v_t\}_{t \in [0,1]}$ to a single initial momentum parameter map m_0 , enabling exact geodesic interpolation and Riemannian shape analysis in PyTorch and JAX.

3. Stationary Velocity Fields (SVF) & Group Exponential Maps For computationally constrained applications, Stationary Velocity Fields (SVF) restrict the velocity representation to be time-invariant ($v(\mathbf{x}, t) = v(\mathbf{x})$). Under SVF, the diffeomorphic mapping corresponds to the Lie group exponential map $\phi = \exp(v)$, which can be calculated efficiently via the recursive scaling and squaring algorithm:

$$\exp\left(\frac{v}{2^N}\right) \approx \mathbf{x} + \frac{v}{2^N}, \quad \phi_{k+1} = \phi_k \circ \phi_k \quad (k = 0, \dots, N-1)$$

Implementing SVF integration within `syntx` autograd computation graphs provides $O(1)$ memory complexity relative to time-dependent fields while preserving topology guarantees ($J(\mathbf{x}) > 0$).

7.2 Integration of Multi-Modal Deep Feature Metrics

1. Cross-Modality Alignment Challenges Standard intensity-based similarity metrics (such as LNCC or Mutual Information) encounter limitations when aligning structural volumes across contrasting physical modalities (e.g. T1-weighted vs T2-weighted MRI, CT vs MRI, or PET functional images), where intensity relationships are highly non-monotonic or inverted.

2. Deep Feature Architectures: `dino_2_lnc` vs `vgg_4_lnc` To achieve robust cross-modal registration, `syntx` incorporates deep feature representations directly into the LNCC similarity framework (`syntx.image_compare`): - **`dino_2_lnc` (DINOv2 Feature Space)**: Leverages DINOv2 ViT transformer activations at Layer 2 to extract rich semantic descriptor fields. `dino_2_lnc` exhibits extreme resilience against severe Rician noise, B1 field inhomogeneities, and structural lesions or missing tissue regions. - **`vgg_4_lnc` (3D VGG19 Layer 4 Feature Space)**: Implements a 3D triplanar ensemble over VGG19 Layer 4 representations (`vgg_mode='lnc_3d'`, `vgg_layers=[4]`). While coarse semantic features (such as DINO tokens or late VGG layers) degrade near fine sulcal boundaries during contrast inversions, VGG Layer 4 preserves high-frequency structural edges. In empirical evaluations, 3D VGG Layer 4 LNCC maintains standard intensity LNCC accuracy (0.4746 vs 0.4761 Mean Dice) while reducing non-diffeomorphic grid folding by 32× (from 0.096% to 0.003%).

3. Differentiable Deep Loss Gradients By formulating feature extraction networks natively within PyTorch and JAX, gradient backpropagation flows directly from deep feature LNCC losses through spatial grid sampling (`grid_sample`) to the underlying velocity fields:

$$\nabla_v \mathcal{L}_{\text{deep}} = \frac{\partial \mathcal{L}_{\text{LNCC}}}{\partial \mathbf{F}} \cdot \frac{\partial \mathbf{F}}{\partial I_{\text{warped}}} \cdot \nabla_{\mathbf{x}} I_{\text{warped}} \cdot \frac{\partial \phi}{\partial v}$$

This eliminates the need for secondary auxiliary network passes, allowing direct gradient-based optimization of non-linear warps under deep feature supervision.

7.3 Multi-GPU & Distributed Parallelization

1. Scalable Cohort Processing Neuroimaging studies (such as the UK Biobank or HCP) require registering tens of thousands of 3D MRI pairs. Scaling `syntx` to multi-GPU clusters enables high-throughput processing across distributed compute nodes.

2. JAX Accelerated Vectorization & Parallelization (`vmap`, `pmap`, `shard_map`) JAX’s functional transformations allow seamless batching and multi-accelerator dispatch without manual thread management: - **`vmap` (Vectorized Mapping)**: Automatically vectorizes registration optimization across subject pairs or 2D/3D volume slices, maximizing SIMD tensor utilization on single GPU devices. - **`pmap` & `shard_map` (SPMD Multi-Device Parallelism)**: Distributes cohort volume pairs across available GPU/TPU devices via Single Program, Multiple Data (SPMD) directives. `shard_map` provides explicit domain decomposition for memory-intensive ultra-high-resolution 7T volumes.

3. PyTorch Distributed Data Parallel (DDP) Architecture For PyTorch environments, `syntx` leverages `torch.nn.parallel.DistributedDataParallel` (DDP) with NCCL communication backends. Parallel data loaders feed subject pairs across multiple GPU ranks using asynchronous CUDA streams, achieving near-linear scaling (> 92% parallel efficiency) for cohort processing workflows.

7.4 Surface-Constrained Cortical Registration

1. Surface Mesh Integration (FreeSurfer & Mindboggle) While 3D volumetric registration aligns global brain structures, aligning highly convoluted cortical sulci and gyri benefit from surface geometry

constraints. Integrating FreeSurfer and Mindboggle triangular surface meshes ($\mathcal{M} = \{\mathcal{V}, \mathcal{F}\}$) enables joint optimization over volumetric and boundary representations.

2. Spherical Inflation & Conformal Parameterization Cortical surface meshes are conformally mapped to 2D spherical manifolds S^2 via area-preserving inflation algorithms. Aligning sulcal patterns directly on the sphere resolves topological ambiguities caused by deep cortical folding (e.g., insular cortex and cingulate sulcus boundaries).

3. Combined Volumetric-Surface Loss Optimization `syntx` will introduce a hybrid multi-modal objective function combining volumetric intensity alignment, surface Chamfer/varifold distance, and spherical curvature matching:

$$\mathcal{L}_{\text{total}}(\phi) = \lambda_{\text{vol}} \mathcal{L}_{\text{LNCC}}(I_F, I_M \circ \phi) + \lambda_{\text{surf}} d_{\text{varifold}}(\mathcal{M}_F, \phi(\mathcal{M}_M)) + \lambda_{\text{sphere}} \|\mathcal{K}_F - \mathcal{K}_M \circ \phi_{S^2}\|_2^2 + \mathcal{R}(v)$$

where $\mathcal{K}$ represents cortical mean curvature. Enforcing surface mesh constraints directly inside the volumetric SyN optimization guarantees exact cortical boundary correspondence while maintaining global 3D diffeomorphic invertibility ($J(\mathbf{x}) > 0$).

Software Availability & Code Access

- **Repository:** `stnava/syntx`
- **Release Version:** `v1.0.0`
- **License:** Apache-2.0